


AWS Databricks Project

Next Create AWS Account, add and verify your credentials

Note: please don't forget to clean-up the resources, once you are done with your project, demo – for personal use. If it is for your organization then resources will be handled by your platform team.



Explore Free Tier products with a new AWS account.

To learn more, visit aws.amazon.com/free.



Sign up for AWS

Root user email address

Used for account recovery and some administrative functions

AWS account name

Choose a name for your account. You can change this name in your account settings after you sign up.

Verify email address

OR

[Sign in to an existing AWS account](#)



Free Tier offers

All AWS accounts can explore 3 different types of free offers, depending on the product used.



Always free
Never expires



12 months free
Start from initial sign-up date



Trials
Start from service activation date

Sign up for AWS

Contact Information

How do you plan to use AWS?

- ☐ Business - for your work, school, or organization
- ☒ Personal - for your own projects

Who should we contact about this account?

Full Name

Abhijit Das

Country Code

+1

Phone Number

[Redacted]

Country or Region

United States



Free Tier offers

All AWS accounts can explore 3 different types of free offers, depending on the product used.



Always free
Never expires



12 months free
Start from initial sign-up date



Trials
Start from service activation date

Sign up for AWS

Contact Information

How do you plan to use AWS?

- ☐ Business - for your work, school, or organization
- ☒ Personal - for your own projects

Who should we contact about this account?

Full Name

Abhijit Das

Country Code Phone Number



Phone Number

1401356009

Country or Region

United States



11:15

Start from service activation date

Country or Region

Address line 1

Address line 2

Apartment, suite, unit, building, floor, etc.

City

State, Province, or Region

California

Postal Code

☒ I have read and agree to the terms of the [AWS Customer Agreement](#).

Agree and Continue (step 2 of 5)

Windows U

Restart to inst
(estimate: 4 m
New updates
and plugged i
your active ho

Secure verification

 We will not charge you for usage below AWS Free Tier limits. We may temporarily hold up to \$1 USD (or an equivalent amount in local currency) as a pending transaction for 3-5 days to verify your identity.



Sign up for AWS

Billing Information

Billing country

Your billing country determines the payment methods available to you to pay for AWS services.

United States ▼

Credit or Debit card number

██



AWS accepts most major credit and debit cards. To learn more about payment options, review our [FAQ](#)

Expiration date

May 2028 ▼

Security code

████████

Cardholder's name

Abhijit Das

Sign up for AWS

Confirm your identity

Before you can use your AWS account, you must verify your phone number. When you continue, the AWS automated system will contact you with a verification code.

How should we send you the verification code?

- ☒ Text message (SMS)
- ☐ Voice call

Country or region code

United States (+1) ▼

Mobile phone number

7148133939

Send SMS (step 4 of 5)

Security Verification



Type the characters as shown above

Reset

Submit

Security check


Please click verify to start your security challenge

Verify



Explore Free Tier products with a new AWS account.

To learn more, visit aws.amazon.com/free.



Sign up for AWS

Root user email address

Used for account recovery and some administrative functions

AWS account name


Choose a name for your account. You can change this name in your account settings after you sign up.

[Verify email address](#)

OR

[Sign in to an existing AWS account](#)

Your AWS Console is ready to use.





Amazon QProductsSolutionsPricingDocumentationLearnPartner NetworkAWS MarketplaceCustomer EnablementEventsExplore More

About AWSContact UsSupportEnglishMy AccountSign In to the Console



Congratulations!

Thank you for signing up with AWS.

We are activating your account, which should take a few minutes. You will receive an email when this is complete.

[Go to the AWS Management Console](#)

[Sign up for another account](#) or [Contact Sales](#)

Hi, I can connect you with an AWS representative or answer questions you have on AWS.

awsSearch[Alt+S]

United States (N. Virginia)aws_databricks

Console Home

Recently visited

- Billing and Cost Management
- IAM Identity Center
- CloudFormation
- S3
- IAM
- EC2

[View all services](#)

Applications (0)

Region: US East (N. Virginia)

us-east-1 (Current Region)Find applications

Name	Description	Region	Originati
The data couldn't be retrieved. Try again later.			

[Go to myApplications](#)

Welcome to AWS

[Getting started with AWS](#)

Learn the fundamentals and find valuable information to get the most out of AWS.

AWS Health

Open issues: 0

Past 7 days

Scheduled changes


Cost and usage

Current month costs: Unable to load


Cost breakdown: Unable to load

Forecasted month end costs: Unable to load

Next step is to open www.databricks.com and click on “Try Databricks”



Select “Use your existing cloud account” to show AWS and Databricks integration and usage.



Enter your details to create Databricks Account with AWS.

Create your Databricks account

1/2

Sign up with your work email to elevate your trial with expert assistance, cloud options and more.

First name

Abhijit

Last name

Das

Email

abhijitaaagdsert2023@gmail.com

Company

Tredence

Title

Sr. Data Engineer

Phone (Optional)

7148133939

Country

United States

What do you want to build and run with Databricks? (Optional)

Learning/Certification, ETL pipelines

By clicking "Continue," you agree to Databricks processing your personal data in accordance with its [Privacy Notice](#).

Continue




Select AWS Cloud Provider



Pick your cloud provider

2/2

You'll need admin access to your cloud account to get started.



Enjoy \$400 in credits during your 14-day AWS trial. Trial ends when credits expire.

By clicking "Continue," you agree to Databricks' [Terms of Service](#).

Continue

Navigate to AWS Marketplace.


Start your free trial on
AWS Marketplace

[Continue to AWS Marketplace](#)



Step 1 – Start your AWS Trial

Click "Try for free" to create a Databricks trial



Step 2 – Set up your account

Follow the instructions to set up your account and



Step 3 – Log in

Once you've created your workspace, click "Launch

aws marketplace

Search

Hello, [User Name]

About Categories Delivery Methods Solutions Resources Your Saved List Become a Channel Partner Sell in AWS Marketplace Amazon Web Services Home Help

[AWS Marketplace](#) > [Data Analytics](#) > [Software as a Service \(SaaS\)](#) > [Databricks Data Intelligence Platform](#)



Databricks Data Intelligence Platform [info](#)

Sold by: [Databricks, Inc.](#)

Free Trial

The Databricks Data Intelligence Platform unlocks the power of data and AI for your entire organization. Enjoy up to \$400 in usage credits during your 14-day free trial. Cancel anytime. After your trial ends, you will automatically be...

[Show more](#)

4.3 ★★★★★ (3) 3 AWS reviews | 404 external reviews

[Overview](#) | [Features](#) | [Pricing](#) | [Legal](#) | [Usage](#) | [Resources](#) | [Support](#) | [Product comparison](#) New | [Reviews](#)

Overview



Highlights

- Simple: Databricks provides a simplified data architecture by unifying data, analytics and AI workloads on one common platform running on Amazon S3.
- Open: Built on top of the world's most successful open source data projects, the Lakehouse Platform unifies your data ecosystem with

aws marketplace

Search

AboutCategoriesDelivery MethodsSolutionsResourcesYour Saved ListBecome a Channel PartnerSell in AWS MarketplaceAmazon W

databricks

Databricks Data Intelligence Platform Info

View purchase options

Sold by: [Databricks, Inc.](#)

OverviewFeaturesPricingLegalUsageResourcesSupportProduct comparison NewReviews

Pricing

Free trial

Try for free

Try this product free according to the free trial terms set by the vendor.

Databricks Data Intelligence Platform Info

View purchase options

Pricing is based on actual usage, with charges varying according to how much you consume. Subscriptions have no end date and may be canceled any time.

Usage costs (1) Info

Dimension	Cost/unit
Databricks Consumption Units	\$1.00

After selecting Try For Free option, it will take you to pricing and usage cost page.

←↻https://aws.amazon.com/marketplace/procurement?productId=85307958-dfde-4d4d-833f-2c8fe68fdddf&offerId=7vv4pt38625pnnr4qnbsv5m1r☆⚙️🔖

Dell(3) Achilles Intro HD...Request #2954846...GmailYouTubeMapsResource classes for...McAfee Security

aws marketplace

Search

AboutCategoriesDelivery MethodsSolutionsResourcesYour Saved ListBecome a Channel PartnerSell in AWS MarketplaceAmazon W

AWS Marketplace > Databricks Data Intelligence Platform > Subscribe to Databricks Data Intelligence Platform

Subscribe to Databricks Data Intelligence Platform Info

To create a subscription, review the pricing information and accept the terms for this software.

Offer details Info

Offer ID

7vv4pt38625pnnr4qnbsv5m1r

Offered by

Databricks, Inc.

Offer type

Public

Pricing and unit details

AWS Marketplace charges for the product based on your usage. Subscriptions have no end date and can be cancelled at any time.

Usage cost (1) Info

We charge additional costs based on your usage that are not included in the contract terms.

Find unit types

< 1 > ⚙️

Units	Cost/unit
Databricks Consumption Units	\$1.00/Units

aws marketplace

Search

AboutCategoriesDelivery MethodsSolutionsResourcesYour Saved ListBecome a Channel PartnerSell in AWS MarketplaceAmazon Web Services

Total amount

Total cost

Product cost is based on usage.

Tax details

Additional taxes may apply

Terms and conditions

Download EULA(s)

By subscribing to this software, you agree to the pricing terms and the seller's [End User License Agreement \(EULA\)](#). You also agree and acknowledge that AWS may, on your behalf, share information about this transaction (including your payment terms) with the respective seller, reseller or underlying provider, as applicable, in accordance with the [AWS Privacy Notice](#). AWS will issue invoices and collect payments from you on behalf of the seller through your AWS account. Your use of AWS services is subject to the [AWS Customer Agreement](#) or other agreement with AWS governing your use of such services. If you are receiving a private offer from a channel partner, you may click [here](#) (for CPPO transaction) or [here](#) (for SPPO transaction) for more information on the channel partner.

Purchase order (PO) number

Info

You can assign unique purchase order numbers to charges to include them on your invoices. [Learn more](#)

Purchase order number options

No purchase order

Add a purchase order

Purchase details

Info

Offer ID

Offered by

Databricks, Inc.

Pricing details

Total charges based on usage

Tax details

Additional taxes may apply

aws marketplace

Search

AboutCategoriesDelivery MethodsSolutionsResourcesYour Saved ListBecome a Channel PartnerSell in AWS MarketplaceAmazon Web Services

Purchase order (PO) number

Info

You can assign unique purchase order numbers to charges to include them on your invoices. [Learn more](#)

Purchase order number options

No purchase order

Add a purchase order

Purchase details

Info

Offer ID

Offered by

Databricks, Inc.

Pricing details

Total charges based on usage

Tax details

Additional taxes may apply

Purchase order numbers

-

Back

Subscribe

aws marketplace

Search

Hello, [redacted]

AboutCategoriesDelivery MethodsSolutionsResourcesYour Saved ListBecome a Channel PartnerSell in AWS MarketplaceAmazon Web Services HomeHelp

Subscribe to Databricks Data Intelligence Platform [Info](#)

To create a subscription, review the pricing information and accept the terms for this software.

Your subscription is in progress. You can now set up your account on the vendor's website.

Set up your account [↗](#)

Offer details [Info](#)

Offer ID

Offered by

Offer type

[redacted]

Databricks, Inc.

Public

Pricing and unit details

AWS Marketplace charges for the product based on your usage. Subscriptions have no end date and can be cancelled at any time.

Usage cost (1) [Info](#)

We charge additional costs based on your usage that are not included in the contract terms.

Find unit types

< 1 > ⚙

Units

Cost/unit

Databricks Consumption Units

\$1.00/Units

aws marketplace

Search

Hello, [redacted]

AboutCategoriesDelivery MethodsSolutionsResourcesYour Saved ListBecome a Channel PartnerSell in AWS MarketplaceAmazon Web Services HomeHelp

[AWS Marketplace](#) > [Databricks Data Intelligence Platform](#) > [Subscribe to Databricks Data Intelligence Platform](#)

Subscribe to Databricks Data Intelligence Platform [Info](#)

To create a subscription, review the pricing information and accept the terms for this software.

✔ Your purchase of Databricks Data Intelligence Platform is complete

Next, set up your account and complete registration on the vendor's website. If you're unable to complete registration, return through the [Manage subscriptions](#) [↗](#) page on AWS Marketplace.

Set up your account [↗](#)

Offer details [Info](#)

Offer ID

Offered by

Offer type

[redacted]

Databricks, Inc.

Public

Pricing and unit details

AWS Marketplace charges for the product based on your usage. Subscriptions have no end date and can be cancelled at any time.

Usage cost (1) [Info](#)

We charge additional costs based on your usage that are not included in the contract terms.

Find unit types

< 1 > ⚙



Units

Cost/unit

Databricks Consumption Units

\$1.00/Units

Now we need to enable AWS Marketplace deployment integration.



aws marketplace

Search

AboutCategoriesDelivery MethodsSolutionsResourcesYour Saved ListBecome a Channel PartnerSell in AWS MarketplaceAmazon V

The deployment parameters integration was successfully enabled.

Step 1

Verify permissions and link Databricks account

Step 2

Configure workspace

Step 3

Review and launch

Verify permissions and link Databricks account

AWS permissions

Review the required permission details for the product below. These [AWS Identity and Access Management \(IAM\)](#) permissions are required to create AWS resources, such as IAM roles and Amazon S3 buckets, for deployment. If you don't have the permissions, copy them and share them with your AWS administrator.

View permissions

Copy permissionsGo to IAM console

Databricks account

We'll use your Databricks account information to configure your product.

Choosing **Create account** will open the Databricks site in a new tab. You must create and link a new Databricks account. After you link the account to your AWS Marketplace subscription, choose **Next**.

Create account

Databricks account

We'll use your Databricks account information to configure your product.

Choosing **Create account** will open the Databricks site in a new tab. You must create and link a new Databricks account. After you link the account to your AWS Marketplace subscription, choose **Next**.

Create account

You must link a Databricks account to continue.

Cancel

Next

Next we need to link AWS Subscription with Databricks Account.

https://accounts.cloud.databricks.com/sign-up?tuuid=4dc4e476-bba5-4b7d-b1e9-ae9b519d04bf

Dell

(3) Achilles Intro HD...

Request #2954846...

Gmail

YouTube

Maps

Resource classes for...

McAfee Security



aws Subscription linking in-progress

Sign up

Databricks 14 day trial


@gmail.com

Abhijit


Das

Sign up

By clicking "Sign up", you agree to the [privacy policy](#).



Now we would be able to see, new Databricks account has been mapped. Click on Next to proceed.



Configure AWS Databricks Workspace

←

↺

https://aws.amazon.com/marketplace/saas/configuration?productId=85307958-dfde-4d4d-833f-2c8fe68fdddf

☆

Dell

(3) Achilles Intro HD...

Request #2954846...

Gmail

YouTube

Maps

Resource classes for...

McAfee Security

aws marketplace

Search

AboutCategoriesDelivery MethodsSolutionsResourcesYour Saved ListBecome a Channel PartnerSell in AWS Marketplace

✔ The deployment parameters integration was successfully enabled.

Databricks Data Intelligence Platform

> Configure and launch

Step 1

[Verify permissions and link Databricks account](#)

Step 2

Configure workspace

Step 3

Review and launch

Configure workspace

Info

Amazon S3 and IAM resources

Info

You can deploy these resources using the default auto-generated names, or provide custom names.

AWS Region for the Databricks workspace

Info

Select a region where the workspace will be created.

US East (N. Virginia)

Workspace assets storage location

Info

Create an Amazon S3 bucket to store your workspace assets like notebooks, libraries, and your default catalog. You will also create a cross-account IAM role that gives Databricks access to the new Amazon S3 bucket. [View permission details](#)

Amazon S3 bucket name for workspace storage

databricks-[REDACTED]cloud-storage-bucket

Bucket name must be unique within the global namespace and follow the bucket naming rules. [See rules for bucket naming](#)

Cross-account IAM role for access to workspace storage

databricks-[REDACTED]cloud-storage-role

Maximum 64 characters. Use alphanumeric and '+-=_.' characters.

←↻🔒https://aws.amazon.com/marketplace/saas/configuration?productId=85307958-dfde-4d4d-833f-2c8fe68fdddf☆

📁 Dell🔴 (3) Achilles Intro HD...📧 Request #2954846...📧 Gmail📺 YouTube📍 Maps🧩 Resource classes for...🌐 McAfee Security

aws marketplace🔍 Search

About ▾Categories ▾Delivery Methods ▾Solutions ▾Resources ▾Your Saved ListBecome a Channel PartnerSell in AWS Marketplace

🟢 The deployment parameters integration was successfully enabled.

Connect an existing Amazon S3 bucket to your new Databricks workspace - *optional* [Info](#)

If you'd like to immediately access an existing Amazon S3 bucket's data in your new Databricks workspace, choose an existing bucket. You'll also create a storage credential IAM role with the necessary permissions to access this bucket. [View permission details](#)

Existing Amazon S3 bucket

Not specified ▾

🔄

VPC resources [Info](#)

We help you deploy the required [Virtual Private Cloud \(VPC\)](#) related resources to launch the Databricks workspace. The following configurations are used for these resources, which you can manage after launch. The [architecture diagram](#) provides an overview of the VPC architecture that will be deployed.

▶ VPC configurations

Acknowledgement

📘 Review the permissions that you are granting Databricks to create and manage resources in your AWS account.

View permissions

aws marketplace🔍 Search

About ▾Categories ▾Delivery Methods ▾Solutions ▾Resources ▾Your Saved ListBecome a Channel PartnerSell in AWS Marketplace

🟢 The deployment parameters integration was successfully enabled.

We help you deploy the required [Virtual Private Cloud \(VPC\)](#) related resources to launch the Databricks workspace. The following configurations are used for these resources, which you can manage after launch. The [architecture diagram](#) provides an overview of the VPC architecture that will be deployed.

▶ VPC configurations

Acknowledgement

📘 Review the permissions that you are granting Databricks to create and manage resources in your AWS account.

View permissions

☒ By checking this box, you acknowledge that you are granting permission to Databricks to create resources in your AWS account, which may incur incremental AWS charges. [Learn more](#)

Cancel

Previous

Next

✓ The deployment parameters integration was successfully enabled.

[Databricks Data Intelligence Platform](#) > Configure and launch

Step 1
[Verify permissions and link
Databricks account](#)

Step 2
[Configure workspace](#)

Step 3
Review and launch

Review and launch [Info](#)

Step 1: Permissions and Databricks authentication

[Review permissions](#)

Databricks account

Status
✓ Linked

Step 2: Workspace

[Edit](#)

Amazon S3 and IAM resources

AWS Region for the Databricks workspace

US East (N. Virginia)

Amazon S3 bucket name for workspace storage

databricks-6xpuuf1pwayycyl1qy2al-cloud-storage-bucket

✓ The deployment parameters integration was successfully enabled.

[Databricks Data Intelligence Platform](#) > Configure and launch

Step 1
[Verify permissions and link
Databricks account](#)

Step 2
[Configure workspace](#)

Step 3
Review and launch

Review and launch [Info](#)

Step 1: Permissions and Databricks authentication

[Review permissions](#)

Databricks account

Status
✓ Linked

Step 2: Workspace

[Edit](#)

Amazon S3 and IAM resources


AWS Region for the Databricks workspace

US East (N. Virginia)


Amazon S3 bucket name for workspace storage

databricks-[REDACTED]-cloud-storage-bucket

To



Then we need to setup VPC configuration (Nodes, VPC, Internet, electric IP, subnets and other network services to prepare and launch the Databricks (Product).



Once you clicked on Launch button, then we will be able see that AWS Databricks service is getting

[illegible]

aws marketplace

Search

Hello, aws_databricks

About Categories Delivery Methods Solutions Resources Your Saved List Become a Channel Partner Sell in AWS Marketplace Amazon Web Services Home

Successfully created AWS resources

To use your new workspace, go to Databricks site. If you experience issues, contact Databricks support.

Open Databricks

Databricks Data Intelligence Platform > Configure and launch

Launch summary

Use the following links to manage your AWS resources. To use your new workspace, go to the Databricks site. If you experience issues with your workspace, contact [Databricks support](#).

Open Databricks

AWS resources

Amazon S3 and IAM resources

Amazon S3 bucket for workspace storage databricks- cloud-storage-bucket	Cross-account IAM role for access to workspace storage databricks- cloud-storage-role	Cross-account IAM role for compute access databricks- cross-account-role
--	--	---

VPC resources

VPC vpc-	Internet gateway igw-	VPC endpoint vpc-	Security group sg-
-----------------------------	--	--------------------------------------	---------------------------------------

Once the deployment is completed successfully, then we can open Databricks Account link.

https://mail.google.com/mail/u/1/#inbox/FMfcgzQZTVtnCVFtRhqswvKxjxkxKZst

Dell (3) Achilles Intro HD... Request #2954846... Gmail YouTube Maps Resource classes for... McAfee Security

Gmail

Search mail

Compose

Inbox 101

Starred

Snoozed

Sent

Drafts

More

Labels +

Thank you for creating a Databricks account. Here are your next steps to get up and running.

Create your workspace

The workspace is where you and your colleagues will do all your data and AI work.


Bring in your data

Connect to cloud storage, upload local files or use one of our integrations with Fivetran, Arcion, Hevo or Rivery to ingest data into your workspace.


You can also use this [step-by-step guide](#) that walks you through creating and setting up the first workspace for your new Databricks account.


Log in to get started

Connect with us by replying to this email if you have any questions.






Default workspace and catalog will be provisioning under the same region. If you would like to create separate catalog, then we need to have separate workspace in different region.






You will receive an email after a successful workspace creation.







Spin-up a Redshift Serverless resource.



Click on Redshift Serverless.



You can get start with default or customized setting



aws

Services

Search

[Alt+S]

Workgroup

Info

Workgroup is a collection of compute resources from which an endpoint is created. Compute properties include network and security settings.

Workgroup name

default-workgroup

Performance and cost controls

Info

Base capacity

Info

128

Network and security

Virtual private cloud (VPC)

vpc-0adcd71e5e10660dd

VPC security group

sg-09d65107196d377be

SSL

Enable

Subnet

subnet-0ef952a304026f4aa,
subnet-00483bfca5fa1dacd,
subnet-0d6de19951777f509,
subnet-0ecfc4104c85d4f33,
subnet-0f84b74453defc115,
subnet-01e82f1aae54ee6e1,

Enhanced VPC routing

Off

Publicly accessible turned off by default. You can edit the Publicly accessible setting in the

Create serverless

It may take a few minutes to complete. After completing the setup, you can work with your data.

Setting up your Amazon Redshift Serverless.

15%

Ease of use with Amazon Redshift Serverless

Access and analyze data without the need to set up, tune, and manage Amazon Redshift clusters.

1

2

3

4

Continue

Create serverless



It may take a few minutes to complete. After completing the setup, you can work with your data.

Setting up your Amazon Redshift Serverless.

✓ Completed



Ease of use with Amazon Redshift Serverless


Access and analyze data without the need to set up, tune, and manage Amazon Redshift clusters.



Continue

The screenshot shows the Amazon Redshift Serverless dashboard. At the top, there's a navigation bar with the AWS logo, 'Services', a search bar, and regional settings. The main header reads 'Amazon Redshift Serverless' and 'Serverless dashboard'. Below this, there's a 'Namespace overview' section with a 'Filter namespace' dropdown set to 'All namespaces'. It displays five metrics, all with a value of 0: 'Total snapshots', 'Datashares in my account', 'Datashares requiring authorization', 'Datashares from other accounts', and 'Datashares requiring association'. Below the metrics is a table titled 'Namespaces / Workgroups' with columns for Namespace, Status, Workgroup, Status, and Average query duration. The table shows 'default-namespace' and 'default-workgroup', both with a status of 'Available'. To the right of the table is a 'Total compute usage' section with a 'Choose a workgroup' dropdown and a 'Last hour' filter. It includes a link to 'AWS Cost Explorer' and a section for 'Total consumed RPU hours'.

Namespace	Status	Workgroup	Status	Average query duration
default-namespace	Available	default-workgroup	Available	



The screenshot shows the 'Create workgroup' configuration page in the Amazon Redshift console. The page is divided into three steps: 'Step 1: Create workgroup', 'Step 2: Choose namespace', and 'Step 3: Review and create'. The 'Create workgroup' step is active. It includes a 'Workgroup' section with a description: 'Workgroup is a collection of compute resources from which an endpoint is created. Compute properties include network and security settings.' Below this is a 'Workgroup name' field with the value 'databricks-demo-wg'. A note states: 'The name must be from 3-64 characters. Valid characters are a-z (lowercase only), 0-9 (numbers), and - (hyphen).' The 'Performance and cost controls' section is also visible, with options for 'Base capacity' (selected) and 'Price-performance target - new'.

Select minimum configuration by creating a new workgroup. We will integrate Redshift Serverless with Databricks to pump a small amount of data from Databricks to Redshift Serverless.

aws

Services

Search

[Alt+S]

United States (N. Virginia)

Review and create

Workflow name

This is a unique name that defines the workflow.

databricks-demo-wg

The name must be from 3-64 characters. Valid characters are a-z (lowercase only), 0-9 (numbers), and - (hyphen).

Performance and cost controls

Info

Set a base capacity to indicate the base amount of Redshift processing units (RPU) that Amazon Redshift can use to run queries. Alternatively, set price-performance target to optimize resources. Amazon Redshift uses AI-driven scaling and optimization to automatically adjust your resources when running queries.

Performance and cost controls

☒ Base capacity

Set the base capacity in Redshift processing units (RPU) used to process your workload.

☐ Price-performance target - new

Choose a price-performance target, and Amazon Redshift will automatically apply AI-driven optimizations to meet your target.

Base capacity

The default value is 128 RPU. To change the base capacity, choose another RPU value.

16

From 8-512 RPU, the value increases in increments of 8 RPU. From 512-1024 RPU, it increases in increments of 32 RPU.

Limits - optional

Set the maximum RPU capacity to limit the compute resources the workflow uses at any point in time.

Limits - optional

Set the maximum RPU capacity to limit the compute resources the workflow uses at any point in time.

Max capacity

Info

To set the maximum capacity in RPU, enter a number.

16

Remove

If base capacity in performance and cost controls is chosen, max capacity must be greater than or equal to your base capacity. This value must be in increments of 8.

Network and security

IP address type

Info

Choose the IP address type for your workflow.

☒ IPv4

Your resources can communicate only over the IPv4 addressing protocol.

☐ Dual-stack mode (New)

Your resources can communicate over IPv4, IPv6, or both. For dual-stack mode, associate an IPv6 CIDR block with a subnet in your VPC using Amazon EC2.

Virtual private cloud (VPC)

This VPC defines the virtual networking environment for this database.

vpc-0adcd71e5e10660dd

Network and security

IP address type [Info](#)

Choose the IP address type for your workgroup.

☒ IPv4

Your resources can communicate only over the IPv4 addressing protocol.

☐ Dual-stack mode [\(New\)](#)

Your resources can communicate over IPv4, IPv6, or both. For dual-stack mode, associate an IPv6 CIDR block with a subnet in your VPC using Amazon EC2.

Virtual private cloud (VPC)

This VPC defines the virtual networking environment for this database.

vpc-0adcd71e5e10660dd ▼

VPC security groups

This VPC security group defines which subnets and IP ranges can be used in the VPC.

Choose one or more security groups ▼

sg-09d65107196d377be ✕

Subnet

The subnet in the chosen VPC that is associated with the specified database.

Choose three or more subnet IDs ▼

subnet-0ef952a304026f4aa ✕

subnet-00483bfca5fa1dacd ✕

subnet-0d6de19951777f509 ✕

subnet-0ecfc4104c85d4f33 ✕

subnet-0f84b74453defc115 ✕

subnet-01e82f1aae54ee6e1 ✕

subnet-0ef952a304026f4aa X	subnet-00483bfca5fa1dacd X
subnet-0d6de19951777f509 X	subnet-0ecfc4104c85d4f33 X
subnet-0f84b74453defc115 X	subnet-01e82f1aae54ee6e1 X

 SSL is now enforced by default for new serverless workgroups.

SSL

SSL is enabled by default to help you get started with your workgroup. You can change the SSL settings after you finish creating your workgroup.

- ☒ **Enable**
Enable SSL for enhanced security
- ☐ **Disable**
Disable SSL

Enhanced VPC routing

Turning on this option routes network traffic between your serverless database and data repositories through a VPC instead of the internet.

- ☐ Turn on enhanced VPC routing

 Publicly accessible turned off by default. You can edit the Publicly accessible setting in the workgroup configuration page after the workgroup is created. [Learn more](#) 

Cancel

Next

Choose namespace

Step 3

Review and create

Namespace is a collection of database objects and users. Data properties include database name and password, permissions, and encryption and security.

- ☒ Create a new namespace
- ☐ Add to an existing namespace

Namespace

This is a unique name that defines the namespace.

databricks-ns

The name must be from 3-64 characters. Valid characters are a-z (lowercase only), 0-9 (numbers), and - (hyphen).

Database name and password

Database name

The name of the first database in the Amazon Redshift Serverless environment.

dev

The name must be 1-64 alphanumeric characters (lowercase only), and it can't be a reserved word.

Admin user credentials

IAM credentials provided as your default admin user credentials. To add a new admin username and password, customize admin user credentials.

- ☒ **Customize admin user credentials**
To use the default IAM credentials, clear this option.

Admin user name

The administrator's user name for the first database.

Customize Admin User Workspace

Database name and password

Database name

The name of the first database in the Amazon Redshift Serverless environment.

dev

The name must be 1-64 alphanumeric characters (lowercase only), and it can't be a reserved word.

Admin user credentials

IAM credentials provided as your default admin user credentials. To add a new admin username and password, customize admin user credentials.

☒ **Customize admin user credentials**

To use the default IAM credentials, clear this option.

Admin user name

The administrator's user name for the first database.

admin

The name must be 1-128 alphanumeric characters, and it can't be a reserved word.

Admin password

Select an option to manage your admin password.

☒ **Manage admin credentials in AWS Secrets Manager** [Info](#)

AWS manages a KMS key that encrypts your data.

Your data is encrypted by default with a key that AWS owns and manages for you. To choose a different key, customize your encryption settings.

☐ **Customize encryption settings (advanced)**

☐ **Generate a password**

Amazon Redshift generates an admin password

admin

The name must be 1-128 alphanumeric characters, and it can't be a reserved word.

Admin password

Select an option to manage your admin password.

☐ **Manage admin credentials in AWS Secrets Manager** [Info](#)

AWS manages a KMS key that encrypts your data.

☐ **Generate a password**

Amazon Redshift generates an admin password.

☒ **Manually add the admin password**

Manually enter the admin password.

Admin user password

The password of the admin user.

test005H123

☒ **Show password**

Must be 8-64 characters long. Must contain at least one uppercase letter, one lowercase letter and one number. Can be any printable ASCII character except "/", "", or "@".

Permissions

 Associate an IAM role so that your serverless endpoint can LOAD and UNLOAD data. You can create an IAM role as the default for this configuration that has the [AmazonRedshiftAllCommandsFullAccess](#)  policy attached. This policy includes permissions to run SQL commands to COPY, UNLOAD, and query data with Amazon Redshift Serverless. This policy also grants permissions to run SELECT statements for related services, such as Amazon S3, Amazon CloudWatch logs, Amazon SageMaker, and AWS Glue. You won't be

Create the default IAM role



 Associate an IAM role so that your serverless endpoint can LOAD and UNLOAD data. You can create an IAM role as the default for this configuration that has the [AmazonRedshiftAllCommandsFullAccess](#)  policy attached. This policy includes permissions to run SQL commands to COPY, UNLOAD, and query data with Amazon Redshift Serverless. This policy also grants permissions to run SELECT statements for related services, such as Amazon S3, Amazon CloudWatch logs, Amazon SageMaker, and AWS Glue. You won't be able to run these SQL commands without an IAM role attached to your namespace.

Specify an S3 bucket for the IAM role to access

To create a new bucket, [visit S3](#) 

☐ No additional S3 bucket

Create the IAM role without specifying S3 buckets.

☒ Any S3 bucket

Allow users that have access to your Redshift Serverless data to also access any S3 bucket and its contents in your AWS account.

☐ Specific S3 buckets

Specify one or more S3 buckets that the IAM role being created has permission to access.

Cancel

Create IAM role as default

Associated IAM roles (1/1)

Set default ▼

Actions ▼

Associate IAM roles

Create IAM role

Create, associate, or remove an IAM role. You can associate up to 50 IAM roles. You can also choose an IAM role and set it as the default.

Find associated iam roles

< 1 > ⚙

☒	IAM roles ↗	▼	Status	Role type
☒	AmazonRedshift-CommandsAccessRole-██████████1...		Not applied	Default

☒	IAM roles ↗	▼	Status	Role type
☒	AmazonRedshift-CommandsAccessRole-20250306T21...		Not applied	Default

Encryption and security

 Your data is encrypted by default with an AWS owned key. To choose a different key, customize your encryption settings.

☐ Customize encryption settings (advanced)

Audit logging [Info](#)

Collects logging information for the database.

Export these logs:

- ☐ User log
- ☐ Connection log
- ☐ User activity log

Cancel

Previous

Next

Step 1

[Create workgroup](#)

Step 2

[Choose namespace](#)

Step 3

Review and create

Review and create

Step 1: Create workgroup

Edit

Workgroup

Workgroup is a collection of compute resources from which an endpoint is created. Compute properties include network and security settings.

Workgroup name

databricks-demo-wg

Performance and cost controls [Info](#)

Set a base capacity to indicate the base amount of Redshift processing units (RPU) that Amazon Redshift can use to run queries. Alternatively, set price-performance target to optimize resources. Amazon Redshift uses AI-driven scaling and optimization to automatically adjust your resources when running queries.

Base capacity

16 RPUs

Limits - optional

Limits - optional

Max capacity

16 RPUs

Network and security

Virtual private cloud (VPC)

[vpc-0adcd71e5e10660dd](#) [🔗](#)

VPC security group

[sg-09d65107196d377be](#) [🔗](#)

SSL

Enable

Subnet

subnet-0ef952a304026f4aa,
subnet-00483bfca5fa1dacd,
subnet-0d6de19951777f509,
subnet-0ecfc4104c85d4f33,
subnet-0f84b74453defc115,
subnet-01e82f1aae54ee6e1,

Step 2: Choose namespace

Edit

Namespace

Namespace is a collection of database objects and users. Data properties include database name and password, permissions, and encryption and security.

Step 2: Choose namespace

[Edit](#)

Namespace

Namespace is a collection of database objects and users. Data properties include database name and password, permissions, and encryption and security.

Target namespace

databricks-ns

Database name and password

Database name

dev

Admin user credentials

admin

Permissions

Default IAM role

arn:aws:iam::051826700310:role/service-
role/AmazonRedshift-CommandsAccessRole-
20250306T215720

IAM roles

arn:aws:iam::051826700310:role/service-
role/AmazonRedshift-CommandsAccessRole-
20250306T215720

Database name and password

Database name
dev

Admin user credentials
admin

Permissions

Default IAM role
arn:aws:iam::051826700310:role/service-
role/AmazonRedshift-CommandsAccessRole-
20250306T215720

IAM roles
arn:aws:iam::051826700310:role/service-
role/AmazonRedshift-CommandsAccessRole-
20250306T215720

Encryption and security

AWS KMS encryption
AWS owned key

Audit logging
Off

Cancel

Previous

Create

aws Services Search [Alt+S] United States (N. Virginia) aws_databricks

Amazon Redshift Serverless

Serverless dashboard [Info](#)

[Refresh](#) [Query data](#) [Create workgroup](#)

Namespace overview [Info](#)

Namespace data from your account

Filter namespace: All namespaces ▼

Total snapshots	Datashares in my account	Datashares requiring authorization	Datashares from other accounts	Datashares requiring association
0	0	0	0	0

Namespaces / Workgroups [Info](#)

Namespace	Status	Workgroup	Status	Average query duration
default-namespace	Available	default-workgroup	Available	
databricks-ns	Available	databricks-demo-wg	Creating	

Total compute usage - new [Info](#)

Choose a workgroup: Last hour

To visualize the costs of your total compute usage, go to [AWS Cost Explorer](#)

Total consumed RPU hours

Retrieving a workgroup's total compute usage uses the workgroup's compute

aws

Services

Search

[Alt+S]

United States (N. Virginia)

Amazon Redshift Serverless

Serverless dashboard

Query data

Create workgroup

Namespace overview

Filter namespace

All namespaces

Namespace data from your account

Total snapshots

0

Datashares in my account

0

Datashares requiring authorization

0

Datashares from other accounts

0

Datashares requiring association

0

Namespaces / Workgroups

Namespace

Status

Workgroup

Status

Average query duration

default-namespace

Available

default-workgroup

Available

Total compute usage - new

databricks-demo-wg

Last hour

default-workgroup

databricks-demo-wg

Total consumed RPU hours

0

aws

Services

Search

[Alt+S]

United States (N. Virginia)

Amazon Redshift Serverless

Serverless dashboard

Query data

Create workgroup

Namespace overview

Filter namespace

All namespaces

Namespace data from your account

Total snapshots

0

Datashares in my account

0

Datashares requiring authorization

0

Datashares from other accounts

0

Datashares requiring association

0

Namespaces / Workgroups

Namespace

Status

Workgroup

Status

Average query duration

default-namespace

Available

default-workgroup

Available

Total compute usage - new

databricks-demo-wg



Last hour

default-workgroup

databricks-demo-wg

Total consumed RPU hours

0



Redshift is a flavor of Postgres

aws Services Search [Alt+S] United States (N. Virginia) aws_databricks

Redshift query editor v2

Editor

Queries

Serverless: default-workgroup

Serverless: databricks-de...

native databases (2)

dev

public

sample_data_dev

tickit

tpcds

tpch

external databases (1)

awsdatacatalog

Untitled 1 x

Run Limit 100 Explain Isolated session Serverless: da... dev

Schedule

1 select * from information_schema.tables limit 10

Result 1 (10)

	table_catalog	table_schema	table_name	table_type	self_referencing_col
☐	dev	pg_internal	redshift_auto_health_che...	BASE TABLE	NULL
☐	dev	pg_auto_copy	copy_job_detail	VIEW	NULL
☐	dev	pg_catalog	pg_default_acl	BASE TABLE	NULL
☐	dev	pg_catalog	pg_conf	BASE TABLE	NULL
☐	dev	pg_catalog	pg_library	BASE TABLE	NULL
☐	dev	pg_catalog	pg_shdepend	BASE TABLE	NULL
☐	dev	pg_catalog	pg_statistic_indicator	BASE TABLE	NULL

Enable public accessibility for the work-group

aws Services Search [Alt+S] United States (N. Virginia) aws_

This workgroup is associated with a namespace

To manage database objects and users, navigate to the namespace that this workgroup is associated with.

databricks-ns

databricks-demo-wg Info

Actions Query data

General information

Workgroup	Date created	Endpoint
databricks-demo-wg	March 06, 2025, 21:59 (UTC-08:00)	databricks-demo-wg.051826700310.us-east-1.redshift-serverless.amazonaws.com:5439/dev
Namespace	Status	JDBC URL
databricks-ns	Available	jdbc:redshift://databricks-demo-wg.051826700310.us-east-1.redshift-serverless.amazonaws.com:5439/dev
Workgroup ARN	Base capacity	ODBC URL
arn:aws:redshift-serverless:us-east-1:051826700310:workgroup/565b8db4-f904-4171-98aa-4b786d0230dc	16 RPU	Driver=(Amazon Redshift (x64)); Server=databricks-demo-wg.051826700310.us-east-1.redshift-serverless.amazonaws.com; Database=dev
Workgroup version	Custom domain name	
1.0.108470	-	
	Patch version	
	Patch 188	

This workgroup is associated with a namespace

To manage database objects and users, navigate to the namespace that this workgroup is associated with.

databricks-ns

databricks-demo-wg

Actions

Query data

General information

Workgroup
databricks-demo-wg

Namespace
databricks-ns

Workgroup ARN
arn:aws:redshift-serverless:us-east-1:workgroup/

Workgroup version
1.0.108470

Date created
March 06, 2025, 21:59 (UTC-08:00)

Status
Available

Base capacity
16 RPUs

Custom domain name
-

Patch version
Patch 188

Endpoint
databricks-demo-wg-us-east-1.redshift-serverless.amazonaws.com:5439/dev

JDBC URL
jdbc:redshift://databricks-demo-wg-us-east-1.redshift-serverless.amazonaws.com:5439/dev

ODBC URL
Driver={Amazon Redshift (x64)}; Server=databricks-demo-wg-us-east-1.redshift-serverless.amazonaws.com; Database=dev

Query and database monitoring

Data access

Limits

Performance

Tags

Network and security

Edit

Virtual private cloud (VPC)
vpc-

VPC endpoint ID
vpce-

VPC security group
sg-

Subnet
subnet-
subnet-
subnet-
subnet-
subnet-
subnet-

SSL
Enabled

Enhanced VPC routing
Off

Publicly accessible
Allow instances and devices outside the VPC to connect to your database through the instance endpoint.

Turned Off

IP address type
IPv4

VPC security groups

This VPC security group defines which subnets and IP ranges can be used in the VPC.

Choose one or more security groups

sg-XXXXXXXXXX X

Subnet

The subnet in the chosen VPC that is associated with the specified database.

Choose three or more subnet IDs

subnet-XXXXXXXXXX X

subnet-XXXXXXXXXX X

subnet-XXXXXXXXXX X

subnet-XXXXXXXXXX X

subnet-XXXXXXXXXX X

subnet-XXXXXXXXXX X

SSL

SSL is enabled by default to help you get started with your workgroup. You can change the SSL settings after you finish creating your workgroup.

☒ Enable

Enable SSL for enhanced security

☐ Disable

Disable SSL

Enhanced VPC routing

Turning on this option routes network traffic between your serverless database and data repositories through a VPC instead of the internet.

☐ Turn on enhanced VPC routing

Publicly accessible

☒ Turn on Publicly accessible

Allow instances and devices outside your VPC to connect to your database through the endpoint.

SSL

SSL is enabled by default to help you get started with your workgroup. You can change the SSL settings after you finish creating your workgroup.

☒ Enable

Enable SSL for enhanced security

☐ Disable

Disable SSL

Enhanced VPC routing

Turning on this option routes network traffic between your serverless database and data repositories through a VPC instead of the internet.

☐ Turn on enhanced VPC routing

Publicly accessible

☒ Turn on Publicly accessible

Allow instances and devices outside your VPC to connect to your database through the endpoint.

 Turning on the Publicly accessible feature grants outside sources access your Redshift Serverless instance. This instance becomes public and outside sources can connect to it. [Learn more](#)

 Your might not be able to connect for about 10 minutes, while the setting changes.

Cancel

Save changes

databricks-demo-wg [Info](#)

[Actions](#)[Query data](#)

General information



Workgroup
databricks-demo-wg

Namespace
databricks-ns

Workgroup ARN
[arn:aws:redshift-serverless:us-east-1:051826700310:workgroup/565b8db4-f904-4171-98aa-4b786d0230dc](#)

Workgroup version
1.0.108470 [Info](#)

Date created
March 06, 2025, 21:59 (UTC-08:00)

Status
Available

Base capacity
16 RPUs

Custom domain name
-

Patch version
[Patch 188](#)

Endpoint
[databricks-demo-wg.051826700310.us-east-1.redshift-serverless.amazonaws.com:5439/dev](#)

JDBC URL
[jdbc:redshift://databricks-demo-wg.051826700310.us-east-1.redshift-serverless.amazonaws.com:5439/dev](#)

ODBC URL
[Driver={Amazon Redshift \(x64\)}; Server=databricks-demo-wg.051826700310.us-east-1.redshift-serverless.amazonaws.com; Database=dev](#)

[Query and database monitoring](#)[Data access](#)[Limits](#)[Performance](#)[Tags](#)

Network and security [Info](#)

[Edit](#)

Virtual private cloud (VPC)
[vpc-0adcd71e5e10660dd](#)

VPC endpoint ID
[vpce-0539b2024ec4a779d](#)

VPC security group
[sg-09d65107196d377be](#)

Subnet
subnet-0ef952a304026f4aa,
subnet-00483bfca5fa1dacd,
subnet-0d6de19951777f509,
subnet-0ecfc4104c85d4f33,
subnet-0f84b74453defc115,
subnet-01e82f1aae54ee6e1,

SSL
Enabled

Enhanced VPC routing
Off

Publicly accessible
Allow instances and devices outside the VPC to connect to your database through the instance endpoint.

Turned On

IP address type
IPv4

Add a Security Group after enabling public accessibility

The screenshot shows the AWS Management Console interface. The top navigation bar includes the AWS logo, a search bar, and various utility icons. The left-hand navigation pane is expanded to show the 'EC2' section, with 'Instances' selected. The main content area displays the 'Security Groups (1)' page. At the top of this page, there are buttons for 'Actions', 'Export security groups to CSV', and 'Create security group'. Below these is a search bar and a 'Clear filters' button. A table lists the security groups, with one entry visible: a security group with ID 'sg-09d65107196d377be', named 'default', associated with VPC 'vpc-0adcd71e5e10660dd'. The table has columns for Name, Security group ID, Security group name, and VPC ID.

aws

Search

[Alt+S]

United States (N. Virginia)

EC2

Dashboard

EC2 Global View

Events

Instances

Instance Types

Launch Templates

Spot Requests

Savings Plans

Reserved Instances

Dedicated Hosts

Capacity Reservations

Images

AMIs

AMI Catalog

Elastic Block Store

Volumes

Security Groups (1/1) Info

Find resources by attribute or tag

sg-09d65107196d377be

Clear filters

1

Actions

Export security groups to CSV

Create security group

Name	Security group ID	Security group name	VPC ID
-	sg-09d65107196d377be	default	vpc-0adcd71e5e10660dd

Inbound rules (1)

Search

1

Manage tags

Edit inbound rules

Name	Security group rule ID	IP version	Type	Protocol	Port range
-	sgr-0c3fbeb5589d829ae	-	All traffic	All	

aws

Search

[Alt+S]

United States (N. Virginia)

EC2 > Security Groups > sg-09d65107196d377be - default > Edit inbound rules

Edit inbound rules Info

Inbound rules control the incoming traffic that's allowed to reach the instance.

Inbound rules Info

Security group rule ID	Type Info	Protocol Info	Port range Info	Source Info	Description - optional Info
sgr-0c3fbeb5589d829ae	All traffic	All	All	Cus...	
-	Redshift	TCP	5439	My IP	

Add rule

sg-

e

/32

Delete

Delete

Cancel Preview changes Save rules

Create a S3 bucket and upload json files

aws

Search

[Alt+S]

United States (N. Virginia)

Amazon S3 > Buckets

Account snapshot - updated every 24 hours

All AWS Regions

View Storage Lens dashboard

General purpose buckets

Directory buckets

General purpose buckets (1)

Info

All AWS Regions

Copy ARN

Empty

Delete

Create bucket

Buckets are containers for data stored in S3.

Find buckets by name

< 1 >

Name	AWS Region	IAM Access Analyzer	Creation date
databricks-6xpuuf1lpwayycy1qy2al-cloud-storage-bucket	US East (N. Virginia) us-east-1	View analyzer for us-east-1	March 4, 2025, 22:45:02 (UTC-08:00)

aws

Search

[Alt+S]

United States (N. Virginia)

Amazon S3 > Buckets > Create bucket

General configuration

AWS Region

US East (N. Virginia) us-east-1

Bucket type

Info

☒ General purpose

Recommended for most use cases and access patterns. General purpose buckets are the original S3 bucket type. They allow a mix of storage classes that redundantly store objects across multiple Availability Zones.

☐ Directory

Recommended for low-latency use cases. These buckets use only the S3 Express One Zone storage class, which provides faster processing of data within a single Availability Zone.

Bucket name

Info

productorder

Bucket names must be 3 to 63 characters and unique within the global namespace. Bucket names must also begin and end with a letter or number. Valid characters are a-z, 0-9, periods (.), and hyphens (-). [Learn More](#)

Copy settings from existing bucket - optional

Only the bucket settings in the following configuration are copied.

Choose bucket

Format: s3://bucket/prefix

Object Ownership

Info

Control ownership of objects written to this bucket from other AWS accounts and the use of access control lists (ACLs). Object ownership determines who can specify access to objects.

☒ ACLs disabled (recommended)

All objects in this bucket are owned by this account. Access to this bucket and its objects is specified using only policies.

☐ ACLs enabled

Objects in this bucket can be owned by other AWS accounts. Access to this bucket and its objects can be specified using ACLs.

Object Ownership [Info](#)

Control ownership of objects written to this bucket from other AWS accounts and the use of access control lists (ACLs). Object ownership determines who can specify access to objects.

☒ **ACLs disabled (recommended)**

All objects in this bucket are owned by this account. Access to this bucket and its objects is specified using only policies.

☐ **ACLs enabled**

Objects in this bucket can be owned by other AWS accounts. Access to this bucket and its objects can be specified using ACLs.

Object Ownership

Bucket owner enforced

Block Public Access settings for this bucket

Public access is granted to buckets and objects through access control lists (ACLs), bucket policies, access point policies, or all. In order to ensure that public access to this bucket and its objects is blocked, turn on Block all public access. These settings apply only to this bucket and its access points. AWS recommends that you turn on Block all public access, but before applying any of these settings, ensure that your applications will work correctly without public access. If you require some level of public access to this bucket or objects within, you can customize the individual settings below to suit your specific storage use cases. [Learn more](#)

☐ **Block all public access**

Turning this setting on is the same as turning on all four settings below. Each of the following settings are independent of one another.

☐ **Block public access to buckets and objects granted through new access control lists (ACLs)**

S3 will block public access permissions applied to newly added buckets or objects, and prevent the creation of new public access ACLs for existing buckets and objects. This setting doesn't change any existing permissions that allow public access to S3 resources using ACLs.

☐ **Block public access to buckets and objects granted through any access control lists (ACLs)**

S3 will ignore all ACLs that grant public access to buckets and objects.

☐ **Block public access to buckets and objects granted through new public bucket or access point policies** **Turning off block all public access might result in this bucket and the objects within becoming public**

AWS recommends that you turn on block all public access, unless public access is required for specific and verified use cases such as static website hosting.

☒ I acknowledge that the current settings might result in this bucket and the objects within becoming public.

Bucket Versioning

Versioning is a means of keeping multiple variants of an object in the same bucket. You can use versioning to preserve, retrieve, and restore every version of every object stored in your Amazon S3 bucket. With versioning, you can easily recover from both unintended user actions and application failures. [Learn more](#)

Bucket Versioning☒ **Disable**☐ **Enable****Tags - optional** (0)

You can use bucket tags to track storage costs and organize buckets. [Learn more](#)

No tags associated with this bucket.

[Add tag](#)

Default encryption [Info](#)

Server-side encryption is automatically applied to new objects stored in this bucket.

- ☒ Server-side encryption with Amazon S3 managed keys (SSE-S3)
- ☐ Server-side encryption with AWS Key Management Service keys (SSE-KMS)
- ☐ Dual-layer server-side encryption with AWS Key Management Service keys (DSSE-KMS)

Using an S3 Bucket Key for SSE-KMS reduces encryption costs by lowering calls to AWS KMS. S3 Bucket Keys aren't supported for DSSE-KMS. [Learn more](#)

- **Advanced settings**

[Cancel](#) [Create bucket](#)

Dell(3) Achilles Intro HD...Request #2954846...GmailYouTubeMapsResource classes for ...McAfee Security

aws

Search[Alt+S]

United States (N. Virginia)

aws_databricks

Amazon S3> Buckets> productorder > Upload

Files and folders (2 total, 7.4 MB)

RemoveAdd filesAdd folder

All files and folders in this table will be uploaded.

Find by name

< 1 >

☐	Name	Folder	Type	Size
☐	part-r-00000-990f5773-9005...	orders/	-	7.1 MB

DestinationInfo

Destination
[s3://productorder](#)

► Destination details

Bucket settings that impact new objects stored in the specified destination.

► Permissions

Grant public access and access to other AWS accounts.

► Properties

Specify storage class, encryption settings, tags, and more.

Cancel

Upload

Upload: status

Close

After you navigate away from this page, the following information is no longer available.

Summary

Destination
s3://productorder

Succeeded
2 files, 7.4 MB (100.00%)

Failed
0 files, 0 B (0%)

Files and folders

Configuration

Files and folders (2 total, 7.4 MB)

Name	Folder	Type	Size	Status	Error
part-r-00000-158b7037-4e...	products/	-	317.4 KB	Succeeded	-
part-r-00000-990f5773-90...	orders/	-	7.1 MB	Succeeded	-

productorder Info

Objects

Metadata

Properties

Permissions

Metrics

Management

Access Points

Objects (6)



Copy S3 URI

Copy URL

Download

Open

Delete

Actions

Create folder

Upload

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Name	Type	Last modified	Size	Storage class
categories/	Folder	-	-	-
customers/	Folder	-	-	-
departments/	Folder	-	-	-
order_items/	Folder	-	-	-
orders/	Folder	-	-	-
products/	Folder	-	-	-

Create / Setup Glue Crawler to read json files

aws

Search

[Alt+S]

United States (N. Virginia)

AWS Glue

Getting started

ETL jobs

Visual ETL

Notebooks

Job run monitoring

Data Catalog tables

Data connections

Workflows (orchestration)

Zero-ETL integrations New

Data Catalog

Databases

Tables

Stream schema registries

Schemas

Connections

Crawlers

Classifiers

Catalog settings

Data Integration and ETL

Legacy pages

Welcome to AWS Glue

Get started by setting up your account and users, cataloging your data, and building ETL jobs to prepare data for analytics.

Prepare your account for AWS Glue

Admins: Grant access to AWS Glue and set a default IAM role.

Set up roles and users

Catalog and search for datasets

View your databases & tables and catalog data using Crawlers.

Go to the Data Catalog

Move and transform data

Updated

Use Zero-ETL integrations to replicate data in near real-time, or ETL jobs to transform data in visual, notebook, or code interface.

Go to Zero-ETL integrations

Go to ETL jobs

Resources and tutorials

Getting started with AWS Glue: [Documentation](#)[AWS Training](#)

Glue in 5 Minutes Videos: [Authoring](#), [GenAI](#), [Monitoring](#), [Orchestration](#)

Data integration and management

Monitor & debug ETL jobs and track usage

Go to job run monitoring

aws

Search

[Alt+S]

United States (N. Virginia)

AWS Glue

Crawlers

AWS Glue

Getting started

ETL jobs

Visual ETL

Notebooks

Job run monitoring

Data Catalog tables

Data connections

Workflows (orchestration)

Zero-ETL integrations New

Data Catalog

Databases

Tables

Stream schema registries

Schemas

Connections

Crawlers

Classifiers

Catalog settings

Crawlers

A crawler connects to a data store, progresses through a prioritized list of classifiers to determine the schema for your data, and then creates metadata tables in your data catalog.

Crawlers (0) [Info](#)

Last updated (UTC)
March 10, 2025 at 05:19:34

Action

Run

Create crawler

View and manage all available crawlers.

Filter crawlers

< 1 >

	Name	State	Schedule	Last run	Last run ti...	Log	Table change...
No resources							
No resources to display.							

AWS Glue

- Getting started
- ETL Jobs
 - Visual ETL
 - Notebooks
 - Job run monitoring
- Data Catalog tables
- Data connections
- Workflows (orchestration)
- Zero-ETL integrations New
- ▼ **Data Catalog**
 - Databases
 - Tables
 - Stream schema registries
 - Schemas
 - Connections
 - Crawlers**
 - Classifiers
 - Catalog settings
- **Data Integration and ETL**
- **Legacy pages**



- Step 1 **Set crawler properties**
- Step 2 Choose data sources and classifiers
- Step 3 Configure security settings
- Step 4 Set output and scheduling
- Step 5 Review and create

Set crawler properties

Crawler details [Info](#)

Name

Name can be up to 255 characters long. Some character set including control characters are prohibited.

Description - optional

Descriptions can be up to 2048 characters long.

▼ Tags - optional

Use tags to organize and identify your resources.

No tags associated with this resource.

[Add new tag](#)

You can add up to 50 more tags. Maximum of 50 tags allowed.

[Cancel](#)[Next](#)

Add data source



Data source

Choose the source of data to be crawled.

S3



Network connection - *optional*

Optionally include a Network connection to use with this S3 target. Note that each crawler is limited to one Network connection so any other S3 targets will also use the same connection (or none, if left blank).



Clear selection

Add new connection

Location of S3 data

- ☒ In this account
- ☐ In a different account

S3 path

Browse for or enter an existing S3 path.

s3://productorder/



View

Browse S3

All folders and files contained in the S3 path are crawled. For example, type s3://MyBucket/MyFolder/ to crawl all objects in MyFolder within MyBucket.

Subsequent crawler runs

This field is a global field that affects all S3 data sources.

- ☒ Crawl all sub-folders
Crawl all folders again with every subsequent crawl.
- ☐ Crawl new sub-folders only
Only Amazon S3 folders that were added since the last crawl will be crawled. If the schemas are compatible, new partitions will be added to existing tables.
- ☐ Crawl based on events
Rely on Amazon S3 events to control what folders to crawl.

Cancel

Add an S3 data source

AWS Glue

Crawlers

>

Add crawler

Step 1

Set crawler properties

Step 2

Choose data sources and classifiers

Step 3

Configure security settings

Step 4

Set output and scheduling

Step 5

Review and create

AWS Glue

Getting started

ETL jobs

Visual ETL

Notebooks

Job run monitoring

Data Catalog tables

Data connections

Workflows (orchestration)

Zero-ETL integrations [New](#)

▼ Data Catalog

Databases

Tables

Stream schema registries

Schemas

Connections

Crawlers

Classifiers

Catalog settings

► Data Integration and ETL

► Legacy pages

Choose data sources and classifiers

Data source configuration

Is your data already mapped to Glue tables?

☒ Not yet

Select one or more data sources to be crawled.

☐ Yes

Select existing tables from your Glue Data Catalog.

Data sources (1) [Info](#)

Edit

Remove

Add a data source

The list of data sources to be scanned by the crawler.

Type	Data source	Parameters
☒ S3	s3://productorder/	Recrawl all

► Custom classifiers - optional

A classifier checks whether a given file is in a format the crawler can handle. If it is, the classifier creates a schema in the form of a StructType object that matches that data format.

Cancel

Previous

Next

AWS

Search

[Alt+S]

United States (N. Virginia)

aws_databricks

AWS Glue

Crawlers

>

Add crawler

Step 1

Set crawler properties

Step 2

Choose data sources and classifiers

Step 3

Configure security settings

Step 4

Set output and scheduling

Step 5

Review and create

Configure security settings

IAM role [Info](#)

Existing IAM role

Choose an IAM role

View

Create new IAM role

Enter new IAM role

AWSGlueServiceRole-Product

Cancel

Create

Checking this box will allow the crawler to use Lake Formation credentials for crawling the data source. If the data source is registered in another account, you must provide the registered account ID. Otherwise, the crawler will crawl only those data sources associated to the account. Only applicable to S3, Glue Catalog, Iceberg, and Hudi data sources.

► Security configuration - optional

Enable at-rest encryption with a security configuration.

Cancel

Previous

Next

AWS Glue

Crawlers

>

Add crawler

AWS Glue

<

Getting started

ETL jobs

Visual ETL

Notebooks

Job run monitoring

Data Catalog tables

Data connections

Workflows (orchestration)

Zero-ETL integrations [New](#)

▼

Data Catalog

Databases

Tables

Stream schema registries

Schemas

Connections

Crawlers

Classifiers

Catalog settings

►

Data Integration and ETL

►

Legacy pages

When the crawler detects schema changes in the data store, how should AWS Glue handle table updates in the data catalog?

☒ Update the table definition in the data catalog

☐ Add new columns only

☐ Ignore the change and don't update the table in the data catalog

New partitions are added, but existing partitions are not updated

☐ Update all new and existing partitions with metadata from the table

Partitions inherit metadata properties — such as their classification, input format, output format, SerDe information, and schema — from their parent table.

How should AWS Glue handle deleted objects in the data store?

☐ Delete tables and partitions from the data catalog

☐ Ignore the change and don't update the table in the data catalog

☒ Mark the table as deprecated in the data catalog

If you run a job that references a deprecated table, the job might fail. Edit jobs that reference deprecated tables to remove them as sources and targets. We recommend that you delete deprecated tables when they are no longer needed.

☒ Create partition indexes automatically

Enable creation of a separate Partition Index for every Catalog table providing efficient lookup for specific partitions.

Crawler schedule

You can define a time-based schedule for your crawlers and jobs in AWS Glue. The definition of these schedules uses the Unix-like [cron](#) syntax. [Learn more](#)

Frequency

On demand

AWS Glue

Crawlers

>

Add crawler

AWS Glue

<

Getting started

ETL jobs

Visual ETL

Notebooks

Job run monitoring

Data Catalog tables

Data connections

Workflows (orchestration)

Zero-ETL integrations [New](#)

►

Data Catalog

Databases

Tables

Stream schema registries

Schemas

Connections

Crawlers

Classifiers

Catalog settings

►

Data Integration and ETL

►

Legacy pages

Set output and scheduling

Step 5

●

Review and create

Step 2: Choose data sources and classifiers

Edit

Data sources (1) [Info](#)

The list of data sources to be scanned by the crawler.

Type	Data source	Parameters
S3	s3://productorder/	Recrawl all

Step 3: Configure security settings

Edit

Configure security settings

IAM role	Security configuration	Lake Formation configuration
AWSGlueServiceRole-Product	-	-

Step 4: Set output and scheduling

Edit

Set output and scheduling

Database	Table prefix - optional	Maximum table threshold - optional	Schedule
product_db	-	-	On demand

Cancel

Previous

Create crawler

Run Glue Crawler

aws

Search

[Alt+S]

United States (N. Virginia)

AWS Glue

Crawlers

Products Crawler

One crawler successfully created

The following crawler is now created: "Products Crawler"

Products Crawler

Last updated (UTC)
March 10, 2025 at 05:34:52

Run crawler

Edit

Delete

Crawler properties

Name	Products Crawler	IAM role	AWSGlueServiceRole-Product	Database	product_db	State	READY
Description	Product Details	Security configuration	-	Lake Formation configuration	-	Table prefix	-
Maximum table threshold	-						

Advanced settings

Crawler runs

Schedule

Data sources

Classifiers

Tags

Crawler runs (0)

The list of crawler runs for this crawler.

Filter data

Filter by a date and time range

Start time (UTC)

End time (UTC)

Current/last duration

Status

DPU hours

Table chan

aws

Search

[Alt+S]

United States (N. Virginia)

AWS Glue

Crawlers

Products Crawler

Starting crawler

Attempting to start run crawler "Products Crawler"

Products Crawler

Last updated (UTC)
March 10, 2025 at 05:34:52

Run crawler

Edit

Delete

Crawler properties

Name	Products Crawler	IAM role	AWSGlueServiceRole-Product	Database	product_db	State	READY
Description	Product Details	Security configuration	-	Lake Formation configuration	-	Table prefix	-
Maximum table threshold	-						

Advanced settings

Create single schema for each S3 path	Inherit schema from table	Table level	Schema updates in the data store
False	False	-	Update the table definition in the data catalog
Object deletion in the data store	Repeat crawls of S3 data stores	Create Partition Index	
Mark the table as deprecated in the data catalog.	Crawl all folders again with every subsequent crawl.	True	

Crawler runs

Schedule

Data sources

Classifiers

Tags

Crawler runs

Schedule

Data sources

Classifiers

Tags

Crawler runs (1)

The list of crawler runs for this crawler.

Filter data

Filter by a date and time range

Start time (UTC)

End time (UTC)

Current/last duration

Status

DPU hours

Table chan

March 10, 2025 at 05:38:06	-	48 s	Running	-	-
----------------------------	---	------	---------	---	---

Crawler successfully starting

The following crawler is now starting: "Products Crawler"

Crawlers

A crawler connects to a data store, progresses through a prioritized list of classifiers to determine the schema for your data, and then creates metadata tables in your data catalog.

Crawlers (1) Info

Last updated (UTC)
March 10, 2025 at 05:39:37

Action

Run

Create crawler

View and manage all available crawlers.

Filter crawlers

< 1 >

☐	Name	State	Schedule	Last run	Last run ti...	Log	Table change...
☐	Products Craw...	Running		-	-	-	-

Crawler successfully starting

The following crawler is now starting: "Products Crawler"

Crawlers

A crawler connects to a data store, progresses through a prioritized list of classifiers to determine the schema for your data, and then creates metadata tables in your data catalog.

Crawlers (1) Info

Last updated (UTC)
March 10, 2025 at 05:40:02

Action

Run

Create crawler

View and manage all available crawlers.

Filter crawlers

< 1 >

☐	Name	State	Schedule	Last run	Last run ti...	Log	Table change...
☐	Products Craw...	Stopping		-	-	-	6 created

Verify Database Creation

Search

[Alt+S]

United States (N. Virginia)

AWS Glue

Databases

Getting started

ETL jobs

Visual ETL

Notebooks

Job run monitoring

Data Catalog tables

Data connections

Workflows (orchestration)

Zero-ETL integrations

Data Catalog

Databases

Tables

Crawler successfully starting

The following crawler is now starting: "Products Crawler"

Databases (1)

A database is a set of associated table definitions, organized into a logical group.

Filter databases

< 1 >

☐	Name	Description	Location URI	Created on (UTC)
☐	product_db	product details	-	March 10, 2025 at 05:30:15

[illegible]

<https://docs.databricks.com/aws/en/data-governance/unity-catalog/hms-federation/hms-federation-glue>

Optional: Create an IAM role with a trust policy that gives Databricks temporary access to AWS Glue resources. (under IAM Roles – will use S3 mount to read source data)

IAM > Roles > Create role

Step 1

Select trusted entity

Step 2

Add permissions

Step 3

Name, review, and create

Select trusted entity

Trusted entity type

☐ AWS service

Allow AWS services like EC2, Lambda, or others to perform actions in this account.

☐ AWS account

Allow entities in other AWS accounts belonging to you or a 3rd party to perform actions in this account.

☐ Web identity

Allows users federated by the specified external web identity provider to assume this role to perform actions in this account.

☐ SAML 2.0 federation

Allow users federated with SAML 2.0 from a corporate directory to perform actions in this account.

☒ Custom trust policy

Create a custom trust policy to enable others to perform actions in this account.

Custom trust policy

Create a custom trust policy to enable others to perform actions in this account.

https://accounts.cloud.databricks.com/?account_id

Dell (3) Achilles Intro HD... Request #2954846... Gmail YouTube Maps Resource classes for... McAfee Security

Account

Welcome to Databricks

Manage your Databricks account

Workspaces

Configure workspace settings. Workspaces contain notebooks, libraries, queries, and workflows

Catalog

Manage metastores as your top-level container for catalogs, schemas (also called databases), views and tables

Usage

View usage details and graphs for your account in Databricks Units (DBU) or estimated costs (in \$USD)

Users & groups

Manage identities for use with jobs, automated tools and systems

Cloud resources

Manage IAM credentials, VPC configurations and S3 bucket permissions for your AWS services

Settings

Configure your Databricks account single sign-on, user provisioning, subscription, and billing

Get started

If you want help setting up your account, please reach out to us at onboarding-help@databricks.com.

Build your team

Not started

aws [Search] [Alt+S] Global

IAM > Roles > Create role

Custom trust policy

Create a custom trust policy to enable others to perform actions in this account.

```
1 {
2   "Version": "2012-10-17",
3   "Statement": [
4     {
5       "Effect": "Allow",
6       "Principal": {
7         "AWS": ["arn:aws:iam::414351767826:role/unity-catalog-prod-UCMasterRole-1455ZJK0TYTL"]
8       },
9       "Action": "sts:AssumeRole",
10      "Condition": {
11        "StringEquals": {
12          "sts:ExternalId": "XXXXXXXXXX"
13        }
14      }
15    ]
16  }
17 }
```

Edit statement Remove

Add actions for STS

Filter actions

☐ All actions (sts:*)

Access level - read

- ☐ GetAccessKeyInfo Info
- ☐ GetCallerIdentity Info
- ☐ GetFederationToken Info
- ☐ GetServiceBearerToken Info
- ☐ GetSessionToken Info

Create policy

aws [Search] [Alt+S]

IAM > Policies > Create policy

Step 1
Specify permissions
Step 2
Review and create

Specify permissions Info

Add permissions by selecting services, actions, resources, and conditions. Build permission statements using the JSON editor.

Policy editor

Visual JSON

```
1 {
2   "Version": "2012-10-17",
3   "Statement": [
4     {
5       "Sid": "GrantCatalogAccessToGlue",
6       "Effect": "Allow",
7       "Action": [
8         "glue:GetDatabase",
9         "glue:GetDatabases",
10        "glue:GetPartition",
11        "glue:GetPartitions",
12        "glue:GetTable",
13        "glue:GetTables",
14        "glue:GetUserDefinedFunction",
15        "glue:GetUserDefinedFunctions",
16        "glue:BatchGetPartition"
17      ],
18       "Resource": ["arn:aws:glue:us-east-1:XXXXXXXXXX:*"]
19     ]
20   ]
21 }
```

Edit statement

Select an example

Step 2 Add permissions

Add permissions [Info](#)

Permissions policies (1041) [Info](#)

Choose one or more policies to attach to your new role.



Filter by Type

Search: dbsglueaccessplcy ✕ All types ▼ 1 match

☐	Policy name ?	Type	Description
☐	dbsglueaccessplcy	Customer managed	-

► Set permissions boundary - optional

Cancel Previous Next



Now we need to modify the trust relationship to allow our account to access this trust policy (complete self-assuming permission).

Copy & paste the ARN to the trust policy.

Identity and Access Management (IAM)

Search IAM

Dashboard

Access management

- User groups
- Users
- Roles**
- Policies
- Identity providers
- Account settings
- Root access management [New](#)

Access reports

- Access Analyzer
- External access
- Unused access
- Analyzer settings
- Credential report

Summary

Creation date: March 09, 2025, 23:41 (UTC-07:00)

ARN: [arn:aws:iam::\[redacted\]:role/dbsglueaccessrole](#)

Link to switch roles in console: [https://signin.aws.amazon.com/saml?role_name=dbsglueaccessrole&account=05187903110](#)

Last activity: -

Maximum session duration: 1 hour

Permissions: Trust relationships | Tags | Last Accessed | Revoke sessions

Trusted entities

Entities that can assume this role under specified conditions.

```

1 {
2   "Version": "2012-10-17",
3   "Statement": [
4     {
5       "Effect": "Allow",
6       "Principal": {
7         "AWS": [
8           "arn:aws:iam::[redacted]:role/unity-catalog-prod-[redacted]",
9           "arn:aws:iam::[redacted]:role/[redacted]dbsglueaccessrole"
10        ]
11      },
12      "Action": "sts:AssumeRole",
13      "Condition": {}
14    }
15  ]
16 }
```

Next, we need to update the policy with this ARN (we haven't update that with the same ARN).

Identity and Access Management (IAM)

Search IAM

Dashboard

Access management

- User groups
- Users
- Roles**
- Policies
- Identity providers
- Account settings
- Root access management [New](#)

Access reports

- Access Analyzer

dbsglueaccesspolicy

Policy name: dbsglueaccesspolicy | Type: Customer managed | Attached entities: 1

[Copy JSON](#) [Edit](#)

```

13 {
14   "glue:GetTables",
15   "glue:GetUserDefinedFunction",
16   "glue:GetUserDefinedFunctions",
17   "glue:BatchGetPartition"
18 },
19 "Resource": [
20   "arn:aws:glue:us-east-1:[redacted]:*"
21 ],
22 },
23 {
24   "Action": [
25     "sts:AssumeRole"
26 ],
27 "Resource": [
28   "arn:aws:iam:[redacted]:role/dbsglueaccessrole"
29 ],
30 "Effect": "Allow"
31 }
32 }
```

Now we have created the IAM role policy for Glue Data Catalog. Next step will be "Open Databricks Workspace -> Catalog -> Quick Access -> External Data (External Location) -> Credentials tab.

databricks

Search data, notebooks, recent, and more... | CTRL + P

Credits workspace

External Data


External Locations | **Credentials** | Connections

Filter credentials | 1 credential


Name	Purpose	Credential Type	Properties	Owner	Comment
workspace	STORAGE	IAM Role	[redacted]	workspace_admins_workspace_28648272...	

[Create credential](#)


Create a Service Credential with the same ARN (used in trust policy – follow [AWS Databricks documentation](#) for more details).




Please navigate to the AWS console and update your IAM role with the custom trust policy provided below to establish a cross-account trust relationship (if you haven't done that)





Then validate the credential configuration






Grant required permission to users in your Organization.



Now login to the Databricks Account Console to create Metastore (for Glue Data Catalog)



Creating Metastore

Copy `aws_region`, `aws_account_id` and `credential` from the following location.

The screenshot shows the Databricks Catalog Explorer interface. The breadcrumb navigation is 'Catalog Explorer > Credentials >'. The selected credential is 'gluedatacataloguceast2demo'. The 'Overview' tab is active, showing a description: 'Glue Data Catalog - Unity Catalog East 2 Region Demo'. Below the description is a table with the following details:

Credential Type	IAM Role
Purpose	SERVICE
IAM role (ARN)	arn:aws:iam::123456789012:role/glue-unity-catalog-east-2-region-demo
External ID	12345678901234567890123456789012

The screenshot shows the AWS Management Console Home page. The 'Recently visited' section lists several services: AWS Glue, IAM, S3, Billing and Cost Management, Amazon Redshift, EC2, IAM Identity Center, and CloudFormation. The 'Applications (0)' section shows no applications are currently listed. The right sidebar contains account information and a 'Sign out' button.

Or we can use Databricks UI to create the Glue Connection (Hive Metastore).

This screenshot is similar to the first one, showing the 'gluedatacataloguceast2demo' credential configuration in the Databricks Catalog Explorer. However, in the top right corner, there is a dropdown menu with the following options: 'Edit', 'Rename', 'Delete', and 'Create AWS Glue connection' (which is highlighted with a yellow icon).

Credentials > gluedatacatalogueest2demo >
Set up connection

- 1 Connection basics
- 2 Connection details
- 3 Catalog basics
- 4 Access
- 5 Metadata

Step 2
Connection details

AWS Region*
The AWS region of targeted Glue catalog.

us-east-2

AWS Account ID*
The AWS account id of targeted Glue catalog.

123456789012

Credential*
The service credential used to access the data source.

On: gluedatacatalogueest2demo

Cancel Back Create connection

Credentials > gluedatacatalogueest2demo >
Set up connection

- 1 Connection basics
- 2 Connection details
- 3 Catalog basics
- 4 Access
- 5 Metadata

Step 1
Connection basics

Connection name*

GlueCatalogConnJC

Connection type*

new Metastore


Metastore type*

AWS Glue


Comment

Cancel Next

Now we will create the Catalog



Successful connect between Databricks and Glue Data Catalog. We should be able to see databases and tables



Create a lakehouse with S3

After creating an S3 bucket, we need to create role and policy to allow Databricks to access S3 bucket for the Lakehouse platform.

<https://docs.databricks.com/aws/en/connect/unity-catalog/cloud-storage/storage-credentials>

The screenshot shows the Amazon S3 console interface. At the top, there's a navigation bar with the Amazon S3 logo and a search bar. Below the navigation bar, the main content area is titled 'Amazon S3' and 'Buckets'. A section titled 'Account snapshot - updated every 24 hours' provides information about storage lens visibility. The 'General purpose buckets' section is active, showing a list of buckets. The 'prod-ux-finder-east2' bucket is highlighted with a yellow box. A callout box points to the 'Copy ARN' button, indicating that the bucket ARN has been copied.

General purpose buckets (1/3) All AWS Regions

Buckets are containers for data stored in S3.

Find buckets by name

Name	AWS Region	IAM Access Analyzer	Creation date
dataanalytics-expusd1pwwyrcytqz1az-cloud-storage-bucket	US East (N. Virginia) us-east-1	View analyzer for us-east-1	March 4, 2025, 22:45:02 (UTC-08:00)
prod-ux-finder	US East (N. Virginia) us-east-1	View analyzer for us-east-1	March 9, 2025, 21:24:33 (UTC-07:00)
prod-ux-finder-east2	US East (Ohio) us-east-2	View analyzer for us-east-2	March 12, 2025, 19:31:41 (UTC-07:00)

Attach new policy

IAM > Roles > Create role

Step 3
Name, review, and create

Role name

Enter a meaningful name to identify this role.

Maximum 64 characters. Use alphanumeric and '-', '_', '@', characters.

Description

Add a short explanation for this role.

Maximum 1000 characters. Use letters (a-z and A-Z), numbers (0-9), hyphens, new lines, or any of the following characters: '-', '@', '<002D>', '<002F>'.

Step 1 Select trusted entity

Step 2 **Add permissions**

Step 3 Name, review, and create

Add permissions info

Permissions policies (1044) info

Choose one or more policies to attach to your new role.

Filter by Type: All types 1 match

Policy name	Type	Description
databricks-lakehouse-demo-pty	Customer managed	

▶ Set permissions boundary - optional

Cancel Previous Next

Identity and Access Management (IAM)

Search IAM

Dashboard

Access management

- User groups
- Users
- Roles**
- Policies
- Identity providers
- Account settings
- Root access management New

Access reports

- Access Analyzer
- External access
- Unused access
- Analyzer settings

databricks-lakehouse-demo-rl info

Summary

Creation date: March 13, 2025, 23:35 (UTC-07:00)

Last activity: -

Maximum session duration: 1 hour

ARN copied: [arn:aws:iam::\[redacted\]:role/databricks-lakehouse-demo-rl](#)

Permissions Trust relationships Tags Last Accessed Revoke sessions

Permissions policies (1) info

You can attach up to 10 managed policies.

Filter by Type: All types

Policy name	Type
databricks-lakehouse-demo-pty	Customer managed

Edit trust policy and add ARN from role.

IAM > Roles > databricks-lakehouse-demo-rl > Edit trust policy

Edit trust policy

```

1 {
2   "Version": "2012-10-17",
3   "Statement": [
4     {
5       "Effect": "Allow",
6       "Principal": {
7         "AWS": [
8           "arn:aws:iam::[redacted]:role/unity-catalog-prod-UCMasterRole-[redacted]",
9           "arn:aws:iam::[redacted]:role/databricks-lakehouse-demo-rl"
10        ],
11      },
12      "Action": "sts:AssumeRole",
13      "Condition": {
14        "StringEquals": {
15          "sts:ExternalId": "[redacted]"
16        }
17      }
18    ]
19  }

```

Now we need to update policy

- Step 1
Modify permissions in databricks-lakehouse-demo-ply
 Step 2
 Review and save

Modify permissions in databricks-lakehouse-demo-ply [info](#)

Add permissions by selecting services, actions, resources, and conditions. Build permission statements using the JSON editor.

Policy editor

```
1 {
2   "Version": "2012-10-17",
3   "Statement": [
4     {
5       "Action": [
6         "s3:GetObject",
7         "s3:PutObject",
8         "s3:DeleteObject",
9         "s3:ListBucket",
10        "s3:GetBucketLocation"
11      ],
12      "Resource": [
13        "arn:aws:s3:::productorder-east2/*",
14        "arn:aws:s3:::productorder-east2"
15      ],
16      "Effect": "Allow"
17    },
18    {
19      "Action": ["sts:AssumeRole"],
20      "Resource": ["arn:aws:iam::[REDACTED]:role/databricks-lakehouse-demo-r1"],
21      "Effect": "Allow"
22    }
23  ]
24 }
```

Now we need to declare on the bucket – about the new role access. – update S3 -> Permission -> Bucket Policy

Amazon S3 > Buckets > productorder-east2 > Edit bucket policy

Amazon S3

- General purpose buckets
 - Directory buckets
 - Table buckets
 - Access Grants
 - Access Points
 - Object Lambda Access Points
 - Multi-Region Access Points
 - Batch Operations
 - IAM Access Analyzer for S3
- Block Public Access settings for this account
- Storage Lens**
 - Dashboards
 - Storage Lens groups
 - AWS Organizations settings
- Feature spotlight
- AWS Marketplace for S3

Bucket policy

The bucket policy, written in JSON, provides access to the objects stored in the bucket. Bucket policies don't apply to objects owned by other accounts. [Learn more](#)

Bucket ARN
 arn:aws:s3:::productorder-east2

Policy

```
1 {
2   "Version": "2012-10-17",
3   "Statement": [
4     {
5       "Sid": "Statement1",
6       "Principal": {
7         "AWS": "arn:aws:iam::[REDACTED]:role/databricks-lakehouse-demo-r1"
8       },
9       "Effect": "Allow",
10      "Action": [
11        "s3:*"
12      ],
13      "Resource": [
14        "arn:aws:s3:::productorder-east2/*",
15        "arn:aws:s3:::productorder-east2"
16      ]
17    }
18  ]
19 }
```

Edit statement
 Statement1 [Remove](#)

Add actions
 Choose a service

Included
 S3

Available
 S3 Express
 S3 Object Lambda
 S3 Outposts
 S3 Tables

Open Databricks Account Console -> Catalog -> Create Metastore (one metastore allow per region):

←

↻

🔒

https://accounts.cloud.databricks.com/data/create?account_id=

📁 Dell🔴 (3) Achilles Intro HD...📧 Request #2954846...📧 Gmail📺 YouTube📍 Maps🗺️ Resource classes for...🛡️ McAfee Security

Account

Workspaces

Catalog

Usage

User management

Cloud resources

Previews

Settings

Catalog > Create metastore >

Create metastore

1 Create metastore

2 Assign to workspaces

Name

databricks-lakehouse-demo-e2

Region

Ohio (us-east-2)

Select the region for your metastore. You will only be able to assign workspaces in this region to this metastore.

S3 bucket path (optional) ⓘ

s3://productorder-east2/

Optional location for storing managed tables data across all catalogs in the metastore. [Learn more](#)

IAM role ARN (optional) ⓘ

arn:aws:iam::[REDACTED]:role/databricks-lakehouse-demo-rl

Enter the IAM role that Databricks will use to access the S3 bucket. [Learn more](#)

Create

Cancel

Configure IAM Role

×

Please navigate to the AWS console and configure your IAM role with the below value as the externalID to establish a cross account trust relationship. [Learn more](#)

S3 bucket path

s3://productorder-east2/

IAM role ARN

arn:aws:iam::[REDACTED]:role/databricks-lakehouse-demo-rl

External Id

[REDACTED]

Skip (not recommended)

IAM role configured

Let's set up your workspace

We're going to send you to your AWS Console to configure your account.

Once you sign in, we'll pre-populate a CloudFormation template that creates an IAM role and S3 bucket for you, then deploys your workspace.

If you encounter any errors during the process, reach out at onboarding-help.databricks.com for troubleshooting guidance.

Workspace Name

databricks-lakehouse-demo-ws-e2

Human readable name for your workspace

AWS Region of the Databricks workspace

us-east-2

AWS Region where the workspace will be created

Cancel

Start Quickstart

To access S3 from Databricks (mount S3 bucket), create access key and secret key with a IAM User (Admin permission)

IAM > Users > databricksdemouser1 > Create access key

Step 1
Access key best practices & alternatives

Step 2 - optional

Set description tag

Step 3

Retrieve access keys

Set description tag - optional

The description for this access key will be attached to this user as a tag and shown alongside the access key.

Description tag value

Describe the purpose of this access key and where it will be used. A good description will help you rotate this access key confidently later.

Databricks User Access Key

Maximum 256 characters. Allowed characters are letters, numbers, spaces representable in UTF-8, and: _ . : / = + - @

Cancel

Previous

Create access key

Use those keys to mount S3 path – Create cluster from “Compute”. Then open a notebook from “workspace” and attached the running cluster.

[Mounting cloud object storage on Databricks | Databricks Documentation](#)

The screenshot shows the Databricks interface. On the left is a sidebar with navigation options like Workspace, Recents, Catalog, Workflows, Compute, Marketplace, SQL, SQL Editor, Queries, Dashboards, Genie, Alerts, Query History, SQL Warehouses, Data Engineering, Job Runs, Data Ingestion, Pipelines, Machine Learning, Playground, Experiments, Features, Models, and Serving. The main area displays a Python notebook titled 'MountS3SourceData'. The code in the notebook mounts an S3 bucket and reads a JSON file. Below the code, a table view shows the data loaded into a DataFrame. The table has columns for path, name, size, and modificationTime.

```

access_key = "..."
secret_key = "..."
encoded_secret_key = secret_key.replace("/", "%2F")
aws_bucket_name = "productorder-east2"
mount_name = "s3productordereast2"
dbutils.fs.mount(f"s3a://{access_key}:{encoded_secret_key}@{aws_bucket_name}", f"/mnt/{mount_name}")
# file_location = f"/mnt/{mount_name}/categories/categories"
# df = spark.read.format("json").option("header", "true").load(file_location)
# df.printSchema()

root
|-- category_department_id: long (nullable = true)
|-- category_id: long (nullable = true)
|-- category_name: string (nullable = true)

```

	path	name	size	modificationTime
1	dbfs/mnt/s3productordereast2/categories/	categories/	0	1742015339804
2	dbfs/mnt/s3productordereast2/customers/	customers/	0	1742015339804
3	dbfs/mnt/s3productordereast2/departments/	departments/	0	1742015339804
4	dbfs/mnt/s3productordereast2/order_items/	order_items/	0	1742015339804
5	dbfs/mnt/s3productordereast2/orders/	orders/	0	1742015339804

Create schemas (bronze, silver and gold) to perform ETL in a medallion architecture.

Create a new schema

A schema is the second layer of Unity Catalog's three-level namespace and organizes tables and views. [Learn more](#)

Schema name*

Storage location

Create a new external location



Location in cloud storage where data for managed tables will be stored. If the storage location is not specified, it will default to catalog location first and then metastore root location if catalog location is not specified as well.

Comment

Cancel


Create

Create Schema with SQL and PySpark



Methods to read data from Source - > Bronze - > Silver - > Gold

Create Job / Workflow to load data from Source to Gold



New

Workspace

Recents

Catalog

Workflows

Compute

Marketplace

SQL

SQL Editor

Queries

Dashboards

Genie

Alerts

Query History

SQL Warehouses

Data Engineering

Job Runs

Data Ingestion

Pipelines

Machine Learning

Playground

Experiments

Features

Search data, notebooks, recents, and more...

CTRL + P

Credits

Workflows > Jobs >

New Job Mar 16, 2025, 02:18 PM ☆

Send feedback

Run now

Runs

Tasks

categories_src_bronze

...2023@gmail.com/ProductRawToSilver

Abhijit Das's Cluster

Task name*

Type*

Source*

Path*

Cluster*

Dependent libraries

Parameters

categories_src_bronze

Notebook

Workspace

/Workspace/Users/abhijitaadiscert2023@gmail.com/ProductRawToSilver

Abhijit Das's Cluster

8 GB - 2 Cores - DBR 15.4 LTS - Spark 3.5.0 - Scala 2.12

+ Add

UI

JSON

Cancel

Create task

Job details

Job ID

Creator

Run as

Tags

Description

Lineage

Schedules & Triggers

Job parameters

Job notifications

None

Add trigger

No job parameters are defined for this job

Edit parameters

No notifications

Workflows > Jobs >

New Job Mar 16, 2025, 02:18 PM ☆

Send feedback

Run now

Runs

Tasks

Runs

Start date

< Previous

Next >

Mar 16

Run total duration

14s

7s

categories_src_bronze

Cancel runs

Start time

Run ID

Launched

Duration

Spark

Status

Error code

Run paramet...

Mar 16, 2025, 0...

11140455...

Manually

19s

Spark UI / Logs

Runni...

Job details

Job ID

Creator

Run as

Tags

Description

Lineage

Git

Schedules & Triggers

Compute

788513563788530

Abhijit Das

Abhijit Das

Add tag

Add description

No lineage information for this job

Learn more

Not configured

Add Git settings

None

Add trigger

Abhijit Das's Cluster

Single node: m5d.large - On-demand - DBR: 15.4 LTS (includes

Workflows > Jobs >

Product-Details-Job ☆

Runs

Tasks

customers_src_bronze

...2023@gmail.com/ProductRawToSilver

Abhijit Das's Cluster

+ Add task

Parameters ⓘ

UIJSON

Key	Value
source_path	productordereast2/customers/ {}
database_name	databricks_lakehouse_demo_ws {}
table_name	customers {}
+ Add	

Notifications ⓘ


+ Add


Retries ⓘ

+ Add


CancelSave task

Below is the complete flow







Gold Table Unity Catalog Lineage




Data load from Bronze -> Silver -> Gold (Table lineage)




Create credential in Redshift, if you haven't created.









Now you are ready to write data from Databricks to Redshift. Writing data from Databricks gold layer to Redshift Serverless.




Install and Integrate PowerBI Desktop with Databricks




Download Power BI Desktop – connection file




Open Power BI Desktop





To open the connection file:








Generate Personal Access Token (PAT) to connect to Databricks from Power BI



Make sure to save the PAT (it won't be visible again). Now we can connect to Databricks from Power BI.



Here is your Power BI report – accessing data from AWS Databricks – gold layer



3M

Sum of category_id

6

Count of department_name

44

Count of customer_state

Average of Order by Customer City

customer_city

Marion

Asheboro

Lodi

Cayey

Toa Baja

Bay Shore

Sheboygan

Roseburg

Indio

Taunton

Guayama

Des Plaines

Piscataway

Everett

Saint Charles

Eugene

Frankfort

Sanford

Woodside

Average of avg_order_value

0

100

200

300

Average Order Details

department_id	department_name	category_id	category_name	customer_fname	customer_lna
2	Fitness	2	Soccer	Alan	Holmes
2	Fitness	2	Soccer	Alice	Smith
2	Fitness	2	Soccer	Alice	Smith
2	Fitness	2	Soccer	Ann	Alexander
2	Fitness	2	Soccer	Ann	James
2	Fitness	2	Soccer	Ann	Smith
2	Fitness	2	Soccer	Ashley	Smith
2	Fitness	2	Soccer	Ashley	Smith
2	Fitness	2	Soccer	Barbara	Cooley
2	Fitness	2	Soccer	Brenda	Bailey
2	Fitness	2	Soccer	Cynthia	Taylor
2	Fitness	2	Soccer	Daniel	Smith
2	Fitness	2	Soccer	David	Smith
2	Fitness	2	Soccer	Deborah	Smith
2	Fitness	2	Soccer	Deborah	Smith
2	Fitness	2	Soccer	Debra	Morales
2	Fitness	2	Soccer	Diana	Cordova
2	Fitness	2	Soccer	Diane	Schmitt
2	Fitness	2	Soccer	Donna	Sims
2	Fitness	2	Soccer	Dorothy	Parrish
2	Fitness	2	Soccer	Douglas	Watts
2	Fitness	2	Soccer	Elizabeth	Ferguson
2	Fitness	2	Soccer	Emma	Smith
2	Fitness	2	Soccer	Frances	Ward
2	Fitness	2	Soccer	Frank	Livingston
2	Fitness	2	Soccer	George	Mason
2	Fitness	2	Soccer	Hannah	Morton

Filters

Visualizations

Data

Page 1

+